



703650 VO Parallel Systems WS2019/2020

OpenMP Basics

Philipp Gschwandtner

Overview

- ▶ main characteristics
- ▶ programming, execution and memory models
- ▶ directives
- ▶ Tales From the Proseminar

Motivation for Using OpenMP

- ▶ OpenMP is one of the easiest parallel programming models & widely available
 - ▶ however, restricted to shared memory and multi-threaded hardware
- ▶ modern hardware encourages use of such models
 - ▶ desktop: AMD 3990X with 64 cores and 128 threads expected 2020...
 - ▶ AMD server: 2x EPYC 77xx with 64 cores / 128 threads = 128/256
 - ▶ Intel server: 8x Xeon Platinum 827x or 828x with 28 cores / 56 threads = 224/448
 - ▶ exotic hardware: SGI Altix UV (“Mach 2” @ JKU in Linz) with 4096/8192
- ▶ also enables access to up to e.g. 64 TB of RAM on SGI Altix UV systems

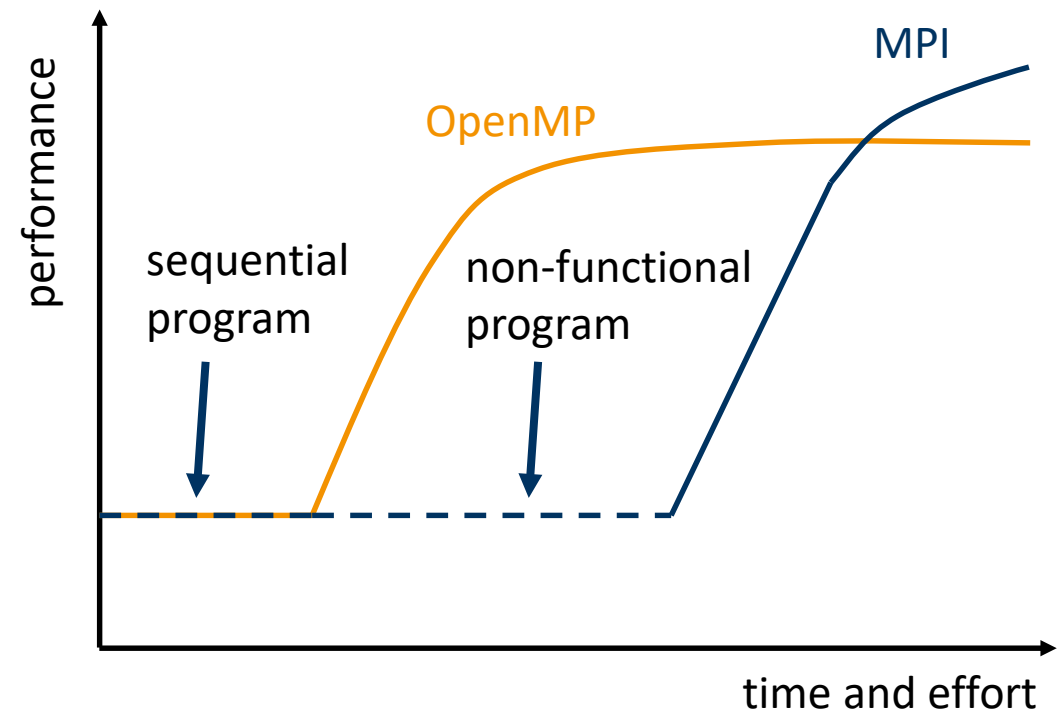
The Advantage of Incremental Parallelization

► MPI

- ▶ initially a sequential program
- ▶ start to parallelize
- ▶ program non-functional until major parts of parallelization present

► OpenMP

- ▶ initially a sequential program
- ▶ parallelize incrementally
- ▶ program remains functional throughout parallelization process



OpenMP

- ▶ thread-based programming model for shared memory parallelism
- ▶ de-facto standard for C/C++ and Fortran
- ▶ maintained by the OpenMP Architecture Review Board
 - ▶ initial release in 1997 (version 1.0 for Fortran)
 - ▶ updates in 1998 (1.0 for C/C++), 2000 (2.0), 2005 (2.5), 2008 (3.0), 2011 (3.1), 2013 (4.0), 2018 (5.0)
 - ▶ this slide set assumes at least OpenMP 3.1!

OpenMP Implementations

- ▶ many implementations available
 - ▶ GCC, LLVM, Intel, Microsoft, etc.
 - ▶ allow to run OpenMP basically on every platform (portability)
 - ▶ that provides decent performance...
 - ▶ compiler and runtime support required
 - ▶ sometimes interchangeable components
 - ▶ check <https://www.openmp.org/resources/openmp-compilers-tools/> for compiler support
- ▶ do not confuse implementation adherence with specification adherence
 - ▶ many minor semantics in OpenMP are implementation-defined

Main Characteristics

- ▶ **compiler-based parallelization model**
 - ▶ tell compiler what code to parallelize and where to put synchronization points
 - ▶ tell compiler whether to share data among threads or create private copies
 - ▶ but **compiler cannot/will not check semantic correctness**
 - ▶ was designed with compiler developers in mind...
- ▶ allows incremental parallelization
- ▶ C/C++ and Fortran bindings
 - ▶ even a research compiler for Java is available...

How to Choose between OpenMP and MPI?

▶ OpenMP

- ▶ language extension, requires compiler support
- ▶ data transfer happens implicitly through shared memory
- ▶ restricted to shared memory parallelism
- ▶ incremental parallelization

▶ MPI

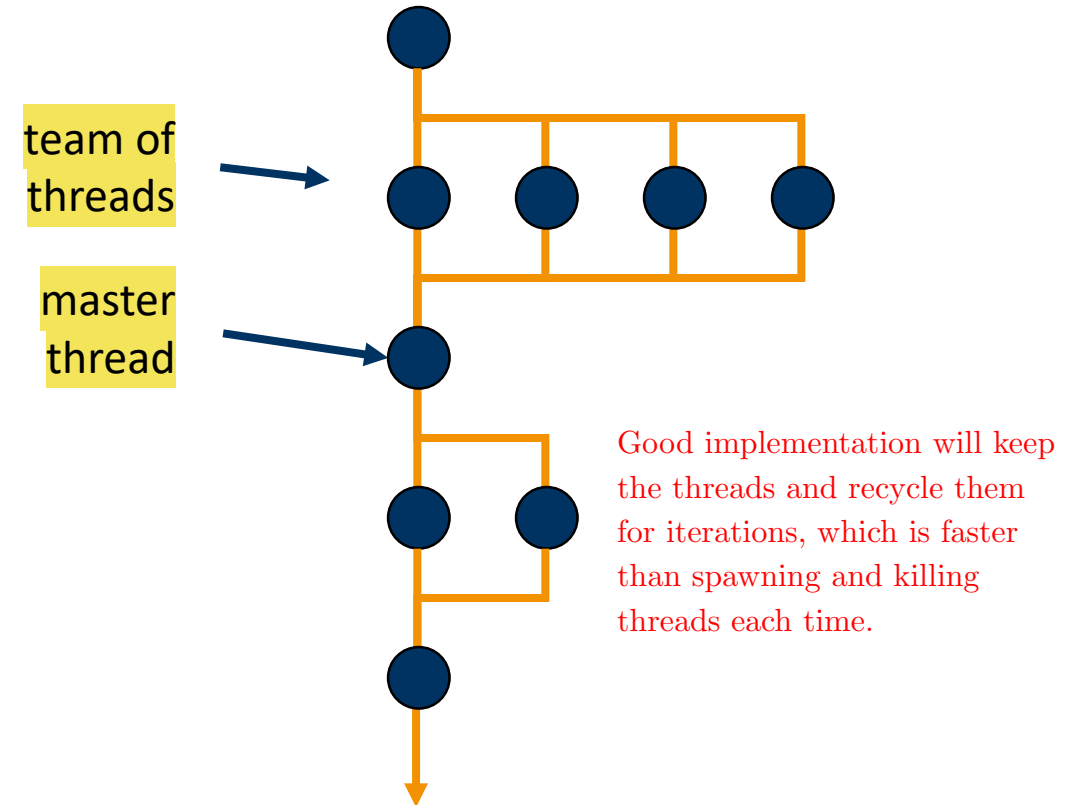
- ▶ library, hence fully compiler-independent
- ▶ data transfer requires explicit message passing
- ▶ supports distributed memory parallelism
- ▶ no incremental parallelization

often: both (*“hybrid”* parallelism)!

Execution Model

▶ fork-join parallelism

- ▶ program starts sequentially
- ▶ parallel regions can be opened, which spawn new threads
- ▶ end of parallel regions synchronize threads
- ▶ afterwards, execution continues sequentially

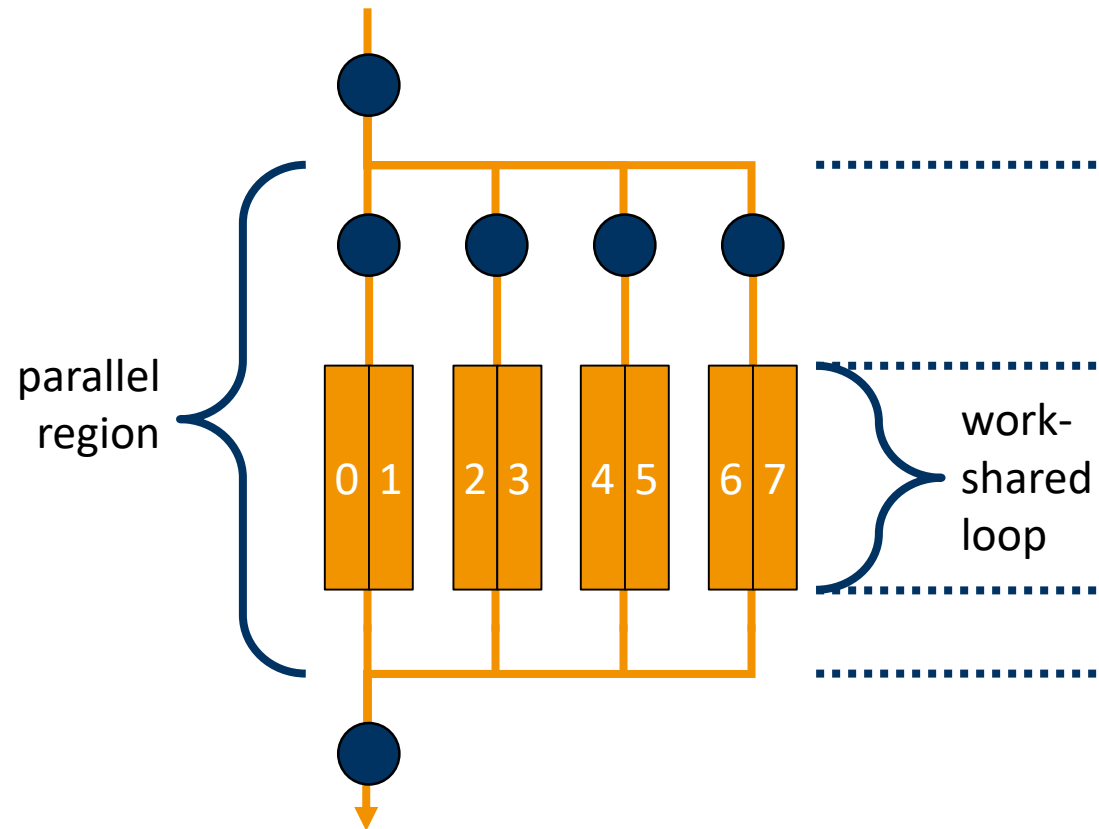


Programming Model

- ▶ mark code regions with **directives** or *pragmas*, e.g.
 - ▶ parallel / sequential regions
 - ▶ work to be distributed
 - ▶ thread synchronization
- ▶ add **clauses** for further information, e.g.
 - ▶ which variables to share, which not to
 - ▶ scheduling strategies
 - ▶ forced execution order
- ▶ any valid OpenMP program must be a valid sequential program if all pragmas are removed!

```
int f = ...  
#pragma omp parallel shared(a,b,c,f)  
    default(none)  
{  
    #pragma omp for  
    for(int i = 0; i < 8; ++i) {  
        c[i] = a[i] + b[i] * f;  
    }  
}
```

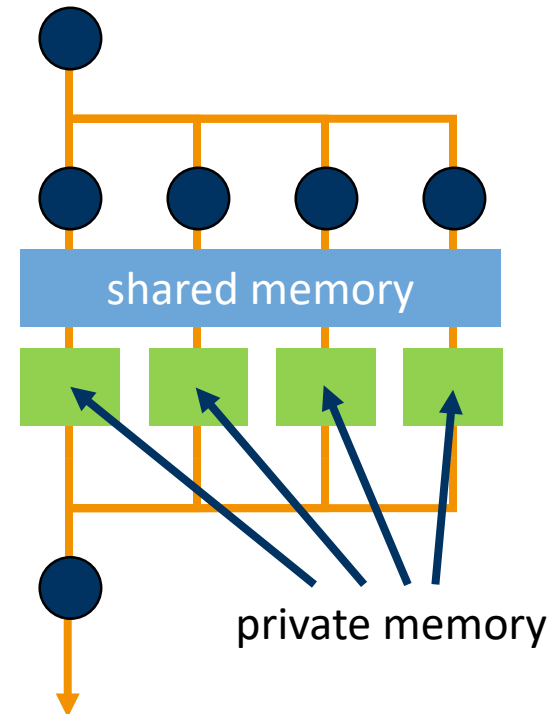
Programming Model cont'd



```
int f = ...  
#pragma omp parallel shared(a,b,c,f)  
    default(none)  
{  
    #pragma omp for  
    for(int i = 0; i < 8; ++i) {  
        c[i] = a[i] + b[i] * f;  
    }  
}
```

Memory Model

- ▶ OpenMP is based on threads
 - ▶ all threads have access to global, shared data
 - ▶ each thread has additional local, private data
 - ▶ modifications to private data are not visible across threads
 - ▶ modifications to shared data are not necessarily immediately visible across threads



Compilation and Execution

- ▶ compile as usual but include OpenMP-specific flag
 - ▶ e.g. `gcc: -fopenmp`
- ▶ execute as usual, but set required environment variables
 - ▶ e.g. `OMP_NUM_THREADS` for controlling degree of parallelism
- ▶ be sure to properly set up your job submission
 - ▶ e.g. for SGE @ LCC2, use `-pe openmp`



OpenMP API



OpenMP API cont'd

▶ **pragmas**

- ▶ control constructs
 - ▶ parallelism & work sharing
- ▶ data sharing
 - ▶ private & shared variables, initialization
- ▶ synchronization
 - ▶ critical & atomic sections, barriers

▶ library functions

- ▶ querying/controlling environment
- ▶ timing
- ▶ locking

▶ environment variables

- ▶ degree and **nesting of parallelism**
- ▶ loop scheduling
- ▶ thread mapping and binding

Pragmas

- ▶ `#pragma omp directive [clause
[, clause]] (newline)`
- ▶ pragmas must be on their own source code line and end with a newline
- ▶ OpenMP directives can often take a number of optional clauses, possibly with parameters
- ▶ pragmas have dynamic and lexical extent
 - ▶ e.g. `#pragma omp for` must always be nested in `#pragma omp parallel`
 - ▶ but not necessarily statically (see example on the right)

```
void bar() {  
    #pragma omp for  
    for(...) { ... }  
}  
  
void foo() {  
    #pragma omp parallel  
    bar();  
}
```


Combined Pragmas

```
int f = ...
#pragma omp parallel shared(a,b,c,f)
    default(none)
{
    #pragma omp for
    for(int i = 0; i < 8; ++i) {
        c[i] = a[i] + b[i] * f;
    }
}
```

```
int f = ...
#pragma omp parallel for shared(a,b,c,f)
    default(none)
for(int i = 0; i < 8; ++i) {
    c[i] = a[i] + b[i] * f;
}
```

Library Functions

▶ querying/controlling environment

- ▶ `omp_get_num_threads()`
- ▶ `omp_get_thread_num()`
- ▶ `omp_get_nested()`
- ▶ `omp_in_parallel()`
- ▶ and a few others, also setters!

▶ timing

- ▶ `omp_get_wtime()`
- ▶ `omp_get_wtick()`

▶ locking

- ▶ `omp_init_lock()`
- ▶ `omp_set_lock()`
- ▶ `omp_unset_lock()`
- ▶ `omp_test_lock()`
- ▶ `omp_destroy_lock()`

Environment Variables

▶ OMP_NUM_THREADS

- ▶ sets the number of threads during execution of parallel regions
- ▶ if dynamic adjustment is enabled, sets the maximum number of threads
- ▶ default is implementation-defined

▶ OMP_SCHEDULE

- ▶ applies to for work sharing only
- ▶ supports `static`, `dynamic`, and `guided` with optional chunk size
- ▶ e.g. `OMP_SCHEDULE=static,4`

▶ OMP_NESTED

- ▶ enables/disables nested parallelism

▶ OMP_PROC_BIND

- ▶ enables/disables thread binding

▶ OMP_PLACES

- ▶ controls mapping threads to hardware resources (e.g. cores or sockets)


parallel Directive

▶ #pragma omp parallel

- ▶ must be followed by a single-entry single-exit statement or an OpenMP construct
- ▶ master thread creates a team of threads, each executing the same code redundantly
- ▶ implicit barrier at the end (threads in team synchronize), only master continues

▶ parallel may also be nested

- ▶ but with great power comes great responsibility...



Hello World in OpenMP

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char** argv) {
    #pragma omp parallel
    {
        int myThreadID = omp_get_thread_num();
        int numThreads = omp_get_num_threads();
        printf("Hello from thread %d of %d\n",
            myThreadID, numThreads);
    }
    return 0;
}
```

```
$ OMP_NUM_THREADS=4 ./hello
Hello from thread 0 of 4
Hello from thread 1 of 4
Hello from thread 3 of 4
Hello from thread 2 of 4
```

Data Sharing Clauses

▶ **private**

- ▶ each thread gets a private copy of variable, no longer “*storage-associated*” with original variable
- ▶ private copy is not initialized
- ▶ often better to declare variables inside parallel region, reduces amount of code (also minimizes “*vertical distance*” in source code)
- ▶ default for variables declared inside parallel region

▶ **shared**

- ▶ each thread references the same, global copy
- ▶ data races if access is not synchronized
- ▶ default for variables declared outside parallel region and global variables, often used for read-only access

▶ **default**

- ▶ can be set to private, shared, or none
- ▶ none helpful for detecting missing variables in clauses (compiler will complain!)

Data Sharing Clauses cont'd

Both will be equal in assembly

```
int f = ...
#pragma omp parallel shared(a,b,c,f)
    default(none)
{
    #pragma omp for
    for(int i = 0; i < 8; ++i) {
        c[i] = a[i] + b[i] * f;
    }
}
```

```
int f = ...
#pragma omp parallel
{
    #pragma omp for
    for(int i = 0; i < 8; ++i) {
        c[i] = a[i] + b[i] * f;
    }
}
```

Data Sharing Clauses cont'd

- ▶ **firstprivate**

- ▶ like `private`, but private copies are initialized with value of copy outside of parallel region

- ▶ **lastprivate**

- ▶ like `private`, but outside copy is set to the private copy of the final iteration (for loops) or last section (sections)

- ▶ **threadprivate**

- ▶ like `private`, but will persist across parallel regions
- ▶ master thread variable is storage-associated with original variable (not the case for `private`!)

reduction Clause

- ▶ performs reduction to a single variable in parallel or loop context
 - ▶ arithmetic ops: +, -, *, max, min
 - ▶ logical ops: &, &&, |, ||, ^
 - ▶ same issues as with MPI regarding associativity of operations!
- ▶ user-defined reductions are possible
 - ▶ need to be declared with #pragma omp declare reduction

```
#pragma omp parallel
{
    #pragma omp for reduction(+:x)
    for(int i = 0; i < 10; ++i) {
        x += i;
    }
}

// or

#pragma omp parallel reduction(-:x)
x -= omp_get_thread_num();
```



Work Sharing



Work Sharing Directives

- ▶ distribute execution of following code region among existing threads of the team
- ▶ must be enclosed in parallel region
- ▶ cannot be directly nested
- ▶ do not launch new threads
- ▶ no barrier on entry
- ▶ implicit barrier on exit
 - ▶ unless `nowait` clause
- ▶ `for`
- ▶ `sections`
- ▶ `single`
- ▶ `task`
- ▶ `simd`

for Directive

- ▶ loop iterations may be executed in parallel
 - ▶ requires loop iterations to be independent
 - ▶ matches SPMD patterns
- ▶ most common form of data parallelism in OpenMP
- ▶ can also take clauses
 - ▶ reduction
 - ▶ schedule
 - ▶ collapse

```
#pragma omp parallel
{
    #pragma omp for
    for(/*init*/; /*test*/; /*inc*/) {
        ...
    }
}
```

for Directive cont'd

- ▶ for loops must have canonical form
 - ▶ defined by the OpenMP specification
 - ▶ requires number of iterations to be known upon loop entry
 - ▶ init, test, and inc expressions must be loop invariant
 - ▶ test only allows <, <=, >, >=
 - ▶ inc only allows common patterns such as ++var, var++, --var, var--, var+=step, var-=step, ...
 - ▶ loop variable must not be written to in loop body
- ▶ C: iterator must be integer or pointer
- ▶ C++: iterator must use random access iterators

```
#pragma omp parallel
{
    #pragma omp for
    for(/*init*/; /*test*/; /*inc*/) {
        ...
    }
}
```

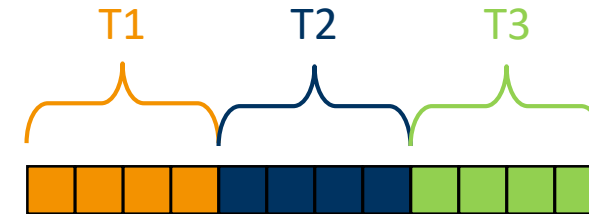
for: schedule Clause

- ▶ specifies method of dividing and assigning chunks
- ▶ static: equally-sized chunks, fixed round-robin assignment
 - ▶ optional: chunk size (default is “approximately equal in size & at most one chunk per thread”)
- ▶ dynamic: equally-sized chunks assigned turn-by-turn
 - ▶ optional: chunk size (default is 1)
- ▶ guided: like dynamic, but chunk size decreases proportionally to no. of unassigned iterations
 - ▶ optional: minimum chunk size (default is 1)
- ▶ also available: auto, runtime

```
#pragma omp parallel
{
    #pragma omp for schedule(dynamic,2)
    for(/*init*/; /*test*/; /*inc*/) {
        ...
    }
}
```

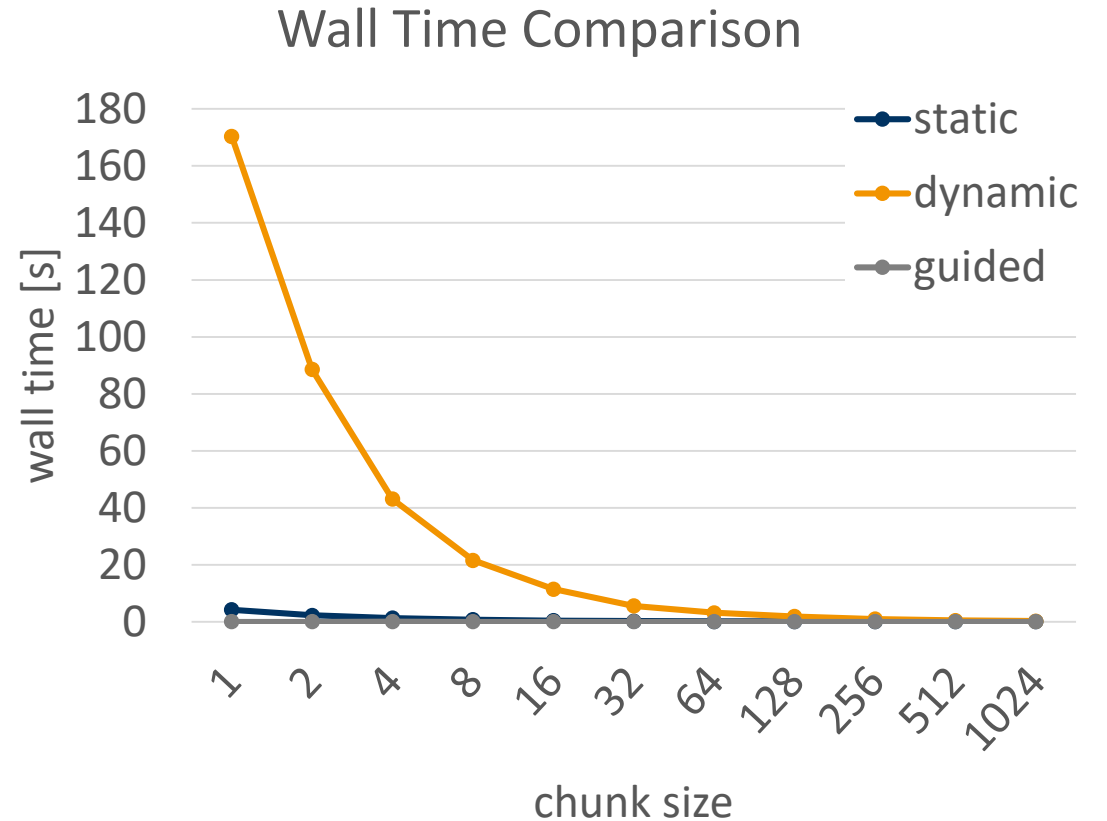
for: schedule Clause

- ▶ specifies method of dividing and assigning chunks
- ▶ static: equally-sized chunks, fixed round-robin assignment
 - ▶ optional: chunk size (default is “approximately equal in size & at most one chunk per thread”)
- ▶ dynamic: equally-sized chunks assigned turn-by-turn
 - ▶ optional: chunk size (default is 1)
- ▶ guided: like dynamic, but chunk size decreases proportionally to no. of unassigned iterations
 - ▶ optional: minimum chunk size (default is 1)
- ▶ also available: auto, runtime



Comparing OpenMP Loop Scheduling Strategies

- ▶ incrementing a 64 bit integer in 10^{11} loop iterations on LCC2, gcc 8.2.0, 8 threads
- ▶ static, dynamic and guided loop scheduling for increasing chunk sizes 2^n with $0 \leq n \leq 10$



for: collapse Clause

- ▶ given a loop nest, the goal is usually to parallelize and distribute outermost loop
 - ▶ minimizes overhead
- ▶ What if outermost loop has low iteration count?
 - ▶ insufficient parallelism for modern systems
- ▶ collapse combines multiple iteration spaces into a single, larger one
 - ▶ allows to exploit more parallelism

```
#pragma omp parallel for collapse(3)
for(int i = 0; i < 3; ++i) {
    for(int j = 0; j < 4; ++j) {
        for(int k = 0; k < 5; ++k) {
            ...
        }
    }
}
```

Careful: You might introduce additional race-conditions this way.

sections Directive

- ▶ sections may be executed concurrently, each by an arbitrary thread of the team
- ▶ matches MPMD patterns
- ▶ easily leads to load imbalance if individual sections not equally work-intensive
 - ▶ also, maximum degree of parallelism limited by number of sections

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        { ... }
        #pragma omp section
        { ... }
        #pragma omp section
        { ... }
    }
}
```

e.g. Asynchronous network calls

single Directive

- ▶ code region will only be executed by a single, arbitrary thread
- ▶ implicit barrier at the end for all threads in the team
- ▶ also available as master variant
 - ▶ like single, but for master thread
 - ▶ no implicit barrier at the end

```
#pragma omp parallel
{
    #pragma omp single
    {
        ...
    }
}
```



Synchronization



Synchronization Directives

- ▶ constructs discussed so far are fairly high-level in terms of synchronization
 - ▶ e.g. unable to enforce specific order or simply wait for all threads
- ▶ OpenMP also offers more fine-grained synchronization directives
 - ▶ barrier
 - ▶ critical
 - ▶ atomic
 - ▶ ordered
 - ▶ flush
 - ▶ master

barrier Directive

- ▶ explicit barrier requested by user
- ▶ threads are not allowed to continue until all have reached the barrier

```
#pragma omp parallel
{
    ...
    #pragma omp barrier
    ...
}
```

critical Directive

- ▶ code region executed by all threads but only one at a time
- ▶ can be **named** to allow finer-grained mutual exclusion
 - ▶ only critical regions with the same name enforce mutual exclusion
 - ▶ all regions without a name have the same name

```
#pragma omp parallel
{
    #pragma omp critical(foo)
    { ... }
    #pragma omp critical(foo)
    { ... }
    #pragma omp critical(bar)
    { ... }
    #pragma omp critical
    { ... }
    #pragma omp critical
    { ... }
}
```

atomic Directive

- ▶ same as `critical`, but restricted to a single memory location and certain operations
- ▶ restriction allows mapping to fast hardware mechanisms
 - ▶ optional specialization clauses `read`, `write`, `update`, and `capture`
- ▶ keeps code hardware- and compiler-independent compared to using intrinsics
 - ▶ but may just be a wrapper for `critical` e.g. when lacking hardware support

```
int count = 0;
#pragma omp parallel
{
    #pragma omp atomic
    count++;
}
```


ordered Directive + Clause

- ▶ enforces critical region with sequential execution order
- ▶ required for strong sequential equivalence
 - ▶ but at high performance cost
- ▶ only efficient if code outside the ordered region is expensive enough
 - ▶ remember guidelines about strong / weak equivalence
 - ▶ check whether your numerical method really needs strong sequential equivalence

```
double total = 0.0;
#pragma omp parallel for ordered
for(int i = 0; i < N; ++i) {
    double part = f(i);
    #pragma omp ordered
    total += part;
}
```

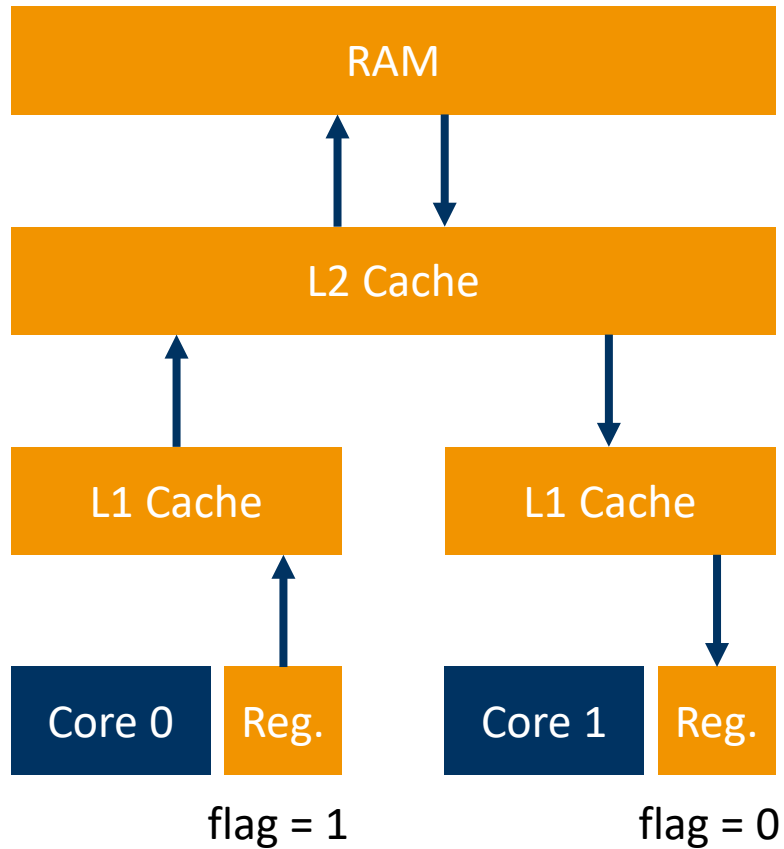
flush Directive

- ▶ problem: writes by one thread are not immediately visible to another in the same parallel region without synchronization
- ▶ example on the right: the second section could hang in while loop
 - ▶ But why?

```
#pragma omp parallel sections
{
    #pragma omp section
    { // producer section
        // produce some data
        flag = 1;
    }
    #pragma omp section
    { // consumer section
        while (flag == 0) { }
        // use data
    }
}
```

flush Directive cont'd

Make sure that data is invalidated if it is shared across threads. (Keep Cachelines in Sync)



```
#pragma omp parallel sections
{
    #pragma omp section
    { // producer section
        // produce some data
        flag = 1;
    }
    #pragma omp section
    { // consumer section
        while (flag == 0) { }
        // use data
    }
}
```

flush Directive cont'd

```
#pragma omp parallel sections
{
    #pragma omp section
    { // producer section
        // produce and flush data
        #pragma omp flush
        #pragma omp atomic write
        flag = 1;
        #pragma omp flush(flag)
    }
}
```

```
#pragma omp section
{ // consumer section
    while (1) {
        #pragma omp flush(flag)
        #pragma omp atomic read
        int temp_flag = flag;
        if(temp_flag == 1) break;
    }
    // use data
}
} // end sections
```

flush Directive cont'd

- ▶ flush is implied at
 - ▶ barrier
 - ▶ entry and exit of critical
 - ▶ exit of
 - ▶ parallel
 - ▶ for
 - ▶ sections
 - ▶ single
 - ▶ set/unset of locks
- ▶ if otherwise required, use flush directive explicitly



Tales From the Proseminar



Tales From the Proseminar

- Which code version is faster? (yes, both have a race condition...)

```
int foo(int dummy) {
    int i, j;
    #pragma omp parallel for private(j)
    for(i = 0; i < 10; ++i) {
        for(j = 0; j < 10; ++j) {
            dummy += i + j;
        }
    }
    return dummy;
}
```

```
int foo(int dummy) {
    #pragma omp parallel for
    for(int i = 0; i < 10; ++i) {
        for(int j = 0; j < 10; ++j) {
            dummy += i + j;
        }
    }
    return dummy;
}
```

Tales From the Proseminar

- ▶ the assembly code is exactly the same
 - ▶ screenshot shows output for gcc 9.2 and -O0 (also works for -O3)
- ▶ <https://godbolt.org/z/SKx8hA>

```
1 foo(int):
2   push    rbp
3   mov     rbp, rsp
4   sub     rsp, 32
5   mov     DWORD PTR [rbp-20], edi
6   mov     eax, DWORD PTR [rbp-20]
7   mov     DWORD PTR [rbp-4], eax
8   lea     rax, [rbp-4]
9   mov     ecx, 0
10  mov     ecx, 0
11  mov     rsi, rax
12  mov     edi, OFFSET FLAT:foo(int) [clone .omp_fn.0]
13  call     GOMP_parallel@plt
14  mov     eax, DWORD PTR [rbp-4]
15  mov     DWORD PTR [rbp-20], eax
16  mov     eax, DWORD PTR [rbp-20]
17  leave
18  ret
19 foo(int) [clone .omp_fn.0]:
20   push    rbp
21   mov     rbp, rsp
22   push    rcx
23   sub     rsp, 40
24   mov     QWORD PTR [rbp-40], rdi
25   call     omp_get_num_threads@plt
26   mov     ebx, eax
27   call     omp_get_thread_num@plt
28   mov     esi, eax
29   mov     eax, 10
30   cqo
31   idiv    ebx
32   mov     ecx, eax
33   mov     eax, 10
34   cqo
35   idiv    ebx
36   mov     eax, edx
37   cmp     esi, eax
38   jl      .L9
39 .L9:
40   imul    esi, ecx
41   mov     edx, esi
42   add     edx, edx
43   lea     ecx, [rax+rcx]
44   cmp     eax, edx
45   jge     .L3
46   mov     DWORD PTR [rbp-20], eax
47 .L7:
48   mov     DWORD PTR [rbp-24], 0
49 .L8:
50   cmp     DWORD PTR [rbp-24], 9
51   jle     .L6
52   add     DWORD PTR [rbp-20], 1
53   cmp     DWORD PTR [rbp-20], edx
54   jl      .L7
55   jmp     .L3
56 .L6:
57   mov     ecx, DWORD PTR [rbp-20]
58   mov     eax, DWORD PTR [rbp-24]
59   add     ecx, eax
60   mov     rax, QWORD PTR [rbp-40]
61   mov     eax, DWORD PTR [rax]
62   add     ecx, eax
63   mov     rax, QWORD PTR [rbp-40]
64   mov     DWORD PTR [rax], ecx
65   add     DWORD PTR [rbp-24], 1
66   jmp     .L8
67 .L4:
68   mov     eax, 0
69   add     ecx, 1
70   jmp     .L9
71 .L3:
72   add     rsp, 40
73   pop     rcx
74   pop     rbp
75   ret
```

```
1 foo(int):
2   push    rbp
3   mov     rbp, rsp
4   sub     rsp, 32
5   mov     DWORD PTR [rbp-20], edi
6   mov     eax, DWORD PTR [rbp-20]
7   mov     DWORD PTR [rbp-4], eax
8   lea     rax, [rbp-4]
9   mov     ecx, 0
10  mov     ecx, 0
11  mov     rsi, rax
12  mov     edi, OFFSET FLAT:foo(int) [clone .omp_fn.0]
13  call     GOMP_parallel@plt
14  mov     eax, DWORD PTR [rbp-4]
15  mov     DWORD PTR [rbp-20], eax
16  mov     eax, DWORD PTR [rbp-20]
17  leave
18  ret
19 foo(int) [clone .omp_fn.0]:
20   push    rbp
21   mov     rbp, rsp
22   push    rcx
23   sub     rsp, 40
24   mov     QWORD PTR [rbp-40], rdi
25   call     omp_get_num_threads@plt
26   mov     ebx, eax
27   call     omp_get_thread_num@plt
28   mov     esi, eax
29   mov     eax, 10
30   cqo
31   idiv    ebx
32   mov     ecx, eax
33   mov     eax, 10
34   cqo
35   idiv    ebx
36   mov     eax, edx
37   cmp     esi, eax
38   jl      .L9
39 .L9:
40   imul    esi, ecx
41   mov     edx, esi
42   add     edx, edx
43   lea     ecx, [rax+rcx]
44   cmp     eax, edx
45   jge     .L3
46   mov     DWORD PTR [rbp-20], eax
47 .L7:
48   mov     DWORD PTR [rbp-24], 0
49 .L8:
50   cmp     DWORD PTR [rbp-24], 9
51   jle     .L6
52   add     DWORD PTR [rbp-20], 1
53   cmp     DWORD PTR [rbp-20], edx
54   jl      .L7
55   jmp     .L3
56 .L6:
57   mov     ecx, DWORD PTR [rbp-20]
58   mov     eax, DWORD PTR [rbp-24]
59   add     ecx, eax
60   mov     rax, QWORD PTR [rbp-40]
61   mov     eax, DWORD PTR [rax]
62   add     ecx, eax
63   mov     rax, QWORD PTR [rbp-40]
64   mov     DWORD PTR [rax], ecx
65   add     DWORD PTR [rbp-24], 1
66   jmp     .L8
67 .L4:
68   mov     eax, 0
69   add     ecx, 1
70   jmp     .L9
71 .L3:
72   add     rsp, 40
73   pop     rcx
74   pop     rbp
75   ret
```


Summary

- ▶ main characteristics
 - ▶ incremental parallelization
- ▶ programming, execution and memory models
 - ▶ based on threads and shared data access
 - ▶ mainly relies on pragmas as programmer interface
- ▶ directives
 - ▶ parallelism, work sharing, data sharing, synchronization
- ▶ Tales From the Proseminar
 - ▶ minimize vertical code distance
 - ▶ specifically: declare for loop iterators in loop header whenever possible